LECTURE:
UNIX ON A REMOTE MACHINE

FIRST STEPS ON A REMOTE MACHINE

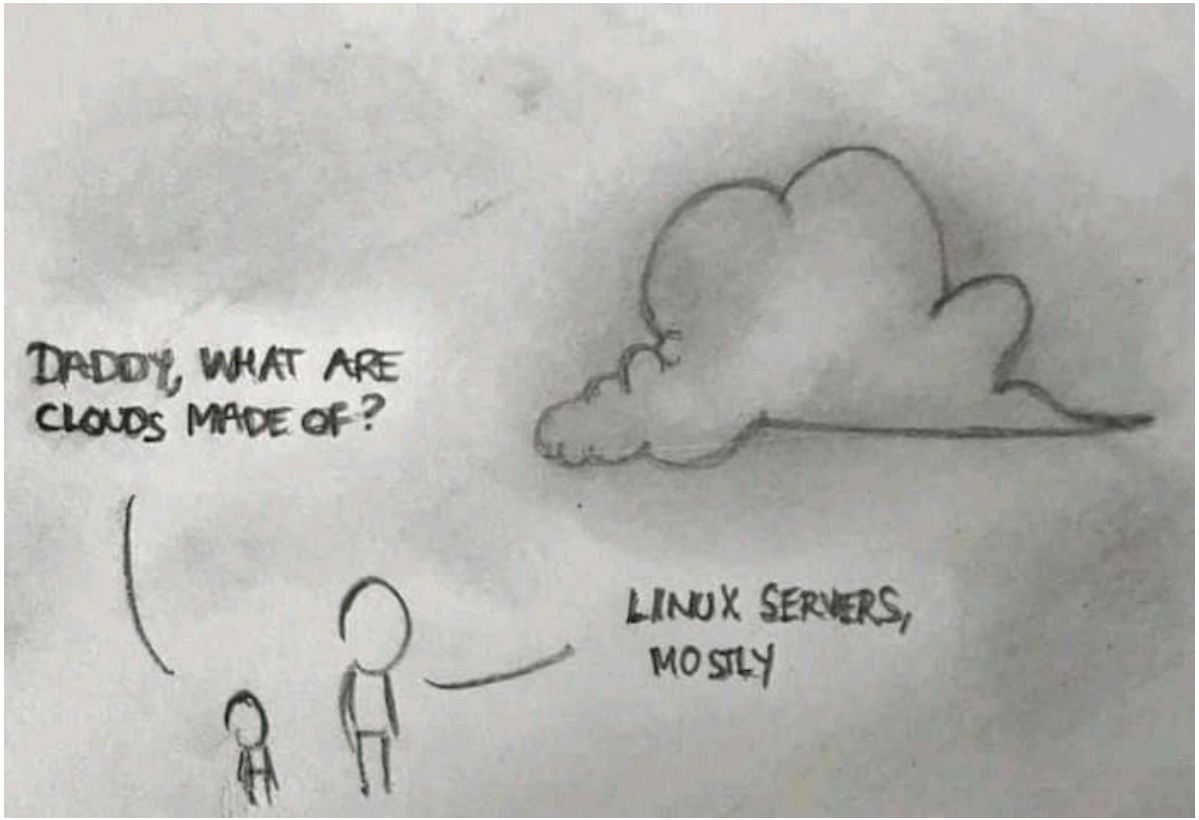

CONNECT VIA SSH & EXECUTING COMMANDS

To `ssh` into a server you execute a command as follows

```
ssh username@remote.machine.name
```

Here we are trying to ssh as user `username` in server `remote.machine.name`. The server can be specified with a URL (like `remote.machine.name`) or an IP (something like username@192.168.1.42).

Executing commands

You run commands directly with `ssh`. `ssh username@server ls` will execute `ls` in the home folder of username. It works with pipes, so

1. `ssh username@server ls | grep PATTERN` will grep locally the remote output of `ls`
2. `ls | ssh username@server grep PATTERN` will grep remotely the local output of `ls`.

SSH KEYS

Key generation

To generate a pair you can run `ssh-keygen`.

```
ssh-keygen -o -a 100 -t rsa -f ~/.ssh/id_rsa
```

You should choose a passphrase, to avoid someone who gets hold of your private key to access authorized servers. Use `ssh-agent` or `gpg-agent` so you do not have to type your passphrase every time.

To check if you have a passphrase and validate it you can run `ssh-keygen -y -f /path/to/key`.

Key based authentication

`ssh` will look into `.ssh/authorized_keys` to determine which clients it should let in. To copy a public key over you can use:

```
cat .ssh/id_rsa.pub | ssh username@remote 'cat >> ~/.ssh/authorized_keys'
```

A simpler solution can be achieved with `ssh-copy-id` where available:

```
ssh-copy-id -i .ssh/id_rsa.pub username@remote
```

COPYING AND DOWNLOADING FILES ON A REMOTE

Copying files over SSH

There are several ways to copy files over ssh:

- `scp` when copying large amounts of files/directories, the secure copy `scp` command is more convenient since it can easily recurse over paths. The syntax is `scp path/to/local_file remote_host:path/to/remote_file`
- `rsync` improves upon `scp` by detecting identical files in local and remote, and preventing copying them again. It also provides more fine grained control over symlinks, permissions and has extra features like the `--partial` flag that can resume from a previously interrupted copy. `rsync` has a similar syntax to `scp`.

Downloading files on a remote machine

Ways to get files on remote machine

- [`wget`] to download using URL or FTP
- [`git clone`] to download a repo from github or gitlab
- [`rclone`] with it you can download datasets from Dropbox, Google Drive, etc

PORT FORWARDING

For example, if we execute

`jupyter notebook` on port `8888`.

We would do `ssh -L 9999:localhost:8888 username@remote_server`

and then navigate to `locahost:9999` in our local machine.

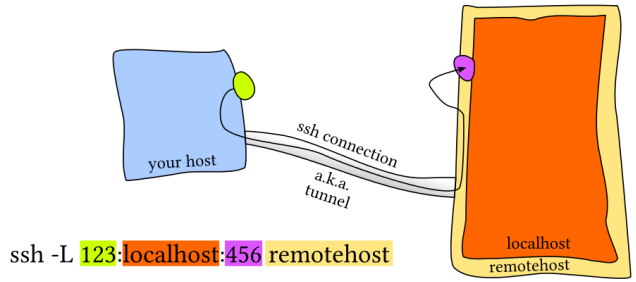

Picture from [this StackOverflow post]

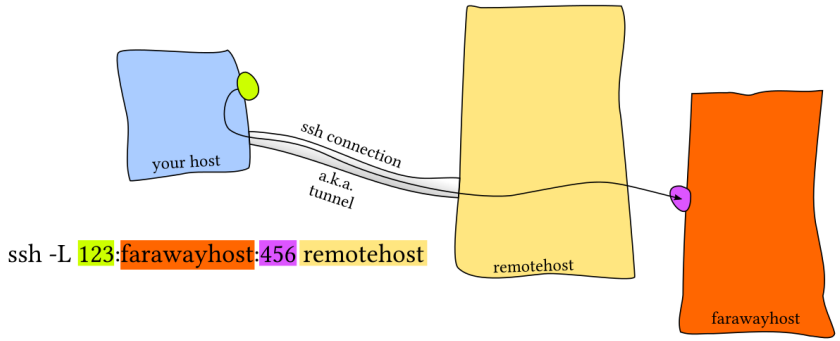

ENVIRONMENT AND CONFIGURATIONS

SSH CONFIGURATION

We have covered many many arguments that we can pass. A tempting alternative is to create shell aliases that look like

```
alias my_server="ssh -i ~/.id_rsa --port 2222 -L 9999:localhost:8888 username@remote_
server
```

However, there is a better alternative using `~/.ssh/config`.

```
Host vm
    User foobar
    HostName 172.16.174.141
    Port 2222
    IdentityFile ~/.ssh/id_rsa
    LocalForward 9999 localhost:8888

# Configs can also take wildcards
Host *.mit.edu
    User foobaz
```

An additional advantage of using the `~/.ssh/config` file over aliases is that other programs like `scp`, `rsync`, `mosh`, etc are able to read it as well and convert the settings into the corresponding flags.

DOTFILES

For `bash`, editing your `.bashrc` or `.bash_profile` will work in most systems. Here you can include commands that you want to run on startup, like the alias we just described or modifications to your `PATH` environment variable. In fact, many programs will ask you to include a line like `export PATH="$PATH:/path/to/program/bin"` in your shell configuration file so their binaries can be found.

Some other examples of tools that can be configured through dotfiles are:

- `bash` - `~/.bashrc`, `~/.bash_profile`
- `git` - `~/.gitconfig`
- `vim` - `~/.vimrc` and the `~/.vim` folder
- `ssh` - `~/.ssh/config`
- `tmux` - `~/.tmux.conf`

Way to learn about customizations is to look through other people's dotfiles: you can find tons of dotfiles repositories on Github --- see the most popular one here (we advise you not to blindly copy configurations though).
Here is another good resource on the topic.

INSPECTING RESOURCES

CPU RAM AND DISK USAGE

- top

```
top - 18:18:35 up 82 days, 30 min,  7 users,  load average: 1.60, 1.12, 0.99
Tasks: 721 total,   1 running, 510 sleeping,   0 stopped,   0 zombie
%Cpu(s):  0.1 us,  0.8 sy,  0.0 ni, 99.1 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem : 26405368+total, 76289200 free, 74934608 used, 11282987+buff/cache
KiB Swap: 26834124+total, 26823603+free,   105216 used. 18677259+avail Mem

  PID USER      PR  NI    VIRT    RES    SHR S  %CPU %MEM     TIME+ COMMAND
18900 root      20   0   12008   2180    864 S   4.3  0.0   2538:10 watch
23286 ebogomo+  20   0   43696   4228   3040 R   1.0  0.0   0:00.18 top
10047 root      20   0       0      0      0 S   0.7  0.0   2174:34 nv_queue
    8 root      20   0       0      0      0 I   0.3  0.0  86:14.60 rcu_sched
 7930 root      20   0   10744   4584   3408 S   0.3  0.0   4:03.70 containerd-shim
10044 root      20   0       0      0      0 S   0.3  0.0   4222:19 nv_queue
10050 root      20   0       0      0      0 S   0.3  0.0   4031:48 nv_queue
17421 root      20   0   23820   2392   1076 S   0.3  0.0 127:39.81 htop
26103 kan       20   0 2780624 203956  14212 S   0.3  0.1  38:34.06 ZMQbg/43
26378 root      20   0   11800   5076   3516 S   0.3  0.0   2:52.40 containerd-shim
    1 root      20   0  225656   9404   6604 S   0.0  0.0   9:02.04 systemd
    2 root      20   0       0      0      0 S   0.0  0.0   0:22.50 kthreadd
    4 root       0 -20       0      0      0 I   0.0  0.0   0:00.00 kworker/0:0H
    6 root       0 -20       0      0      0 I   0.0  0.0   0:00.00 mm_percpu_wq
    7 root      20   0       0      0      0 S   0.0  0.0   0:48.77 ksoftirqd/0
    9 root      20   0       0      0      0 I   0.0  0.0   0:00.00 rcu_bh
```

- htop

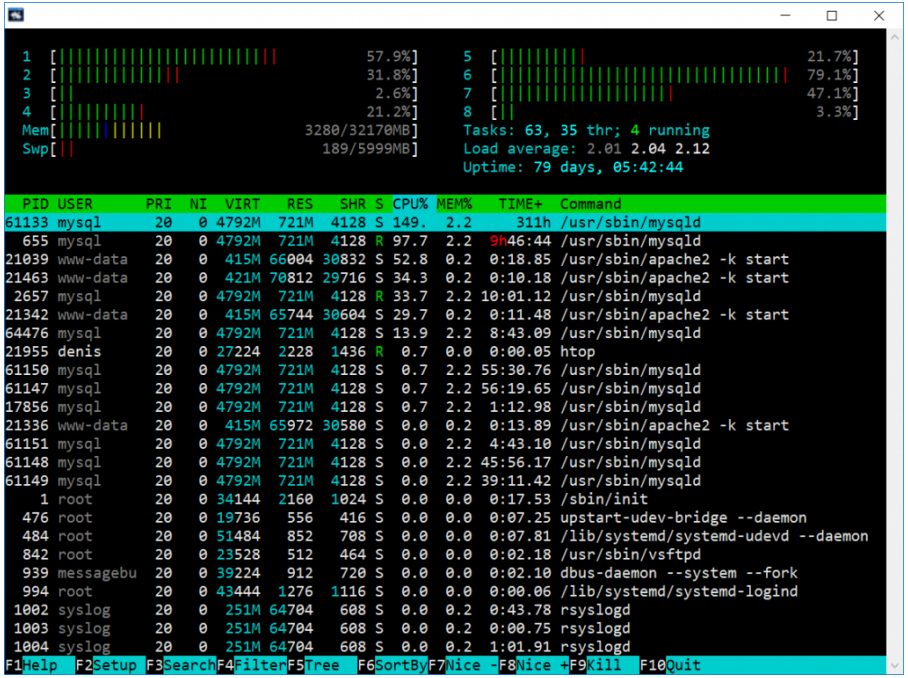

GPU USAGE

- nvidia-smi

```
Wed Sep  9 18:18:49 2020
+-----------------------------------------------------------------------------+
| NVIDIA-SMI 440.64       Driver Version: 440.64       CUDA Version: 10.2     |
|-------------------------------+----------------------+----------------------+
| GPU  Name        Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|         Memory-Usage | GPU-Util  Compute M. |
|===============================+======================+======================|
|   0  GeForce GTX 108...  Off  | 00000000:05:00.0 Off |                  N/A |
|  0%   30C    P8     9W / 280W |     10MiB / 11177MiB |      0%      Default |
+-------------------------------+----------------------+----------------------+
|   1  GeForce GTX 108...  Off  | 00000000:06:00.0 Off |                  N/A |
|  0%   28C    P8     8W / 280W |     10MiB / 11178MiB |      0%      Default |
+-------------------------------+----------------------+----------------------+
|   2  GeForce GTX 108...  Off  | 00000000:09:00.0 Off |                  N/A |
|  0%   28C    P8    12W / 280W |    944MiB / 11178MiB |      0%      Default |
+-------------------------------+----------------------+----------------------+
|   3  GeForce GTX 108...  Off  | 00000000:0A:00.0 Off |                  N/A |
|  0%   28C    P8     9W / 280W |    986MiB / 11178MiB |      0%      Default |
+-------------------------------+----------------------+----------------------+

+-----------------------------------------------------------------------------+
| Processes:                                                       GPU Memory |
|  GPU       PID   Type   Process name                             Usage      |
|=============================================================================|
|    2     21464      C   /home/vikasn/anaconda3/bin/python           429MiB |
|    2     28973      C   /opt/conda/bin/python                       505MiB |
|    3     21464      C   /home/vikasn/anaconda3/bin/python           429MiB |
|    3     28973      C   /opt/conda/bin/python                       547MiB |
+-----------------------------------------------------------------------------+
```

DOCKER:
HOW TO USE A CONTAINER AS A SANDBOX

Skoltech
Skolkovo Institute of Science and Technology

CREATE A SANDBOX

- Docker provides Containers
- A Container is a standard unit of software that packages up code and all its dependencies so the application runs quickly and reliably from one computing environment to another.

- Example of sandbox creation command

```
docker run --rm -it --name fse/student-n \
-v /home/fsestudent/student-n:/root \
--memory 24g --memory-swap 24g \
--cpuset-cpus '0-3' \
--gpus '"device=0,1"' \
nvidia/cuda:8.0-cudnn7-devel-ubuntu16.04
```

WRITE A CODE: TEXT EDITORS

Skoltech
Skolkovo Institute of Science and Technology

VIM, EMACS, NANO

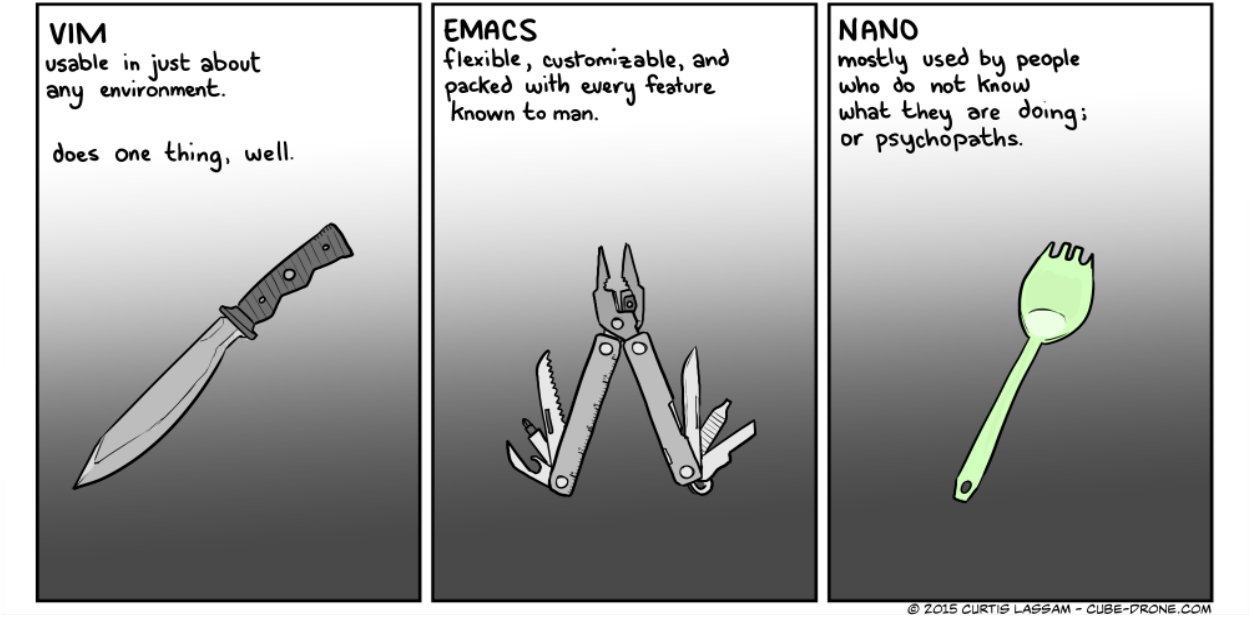

- Links on the tutorials
  - Vim Tutorial
  - Nano Tutorial
  - EMACS Tutorial

JUPYTER, PYCHARM, VSCODE

- Set it up on the remote
- Use it for the coarse file management on your remote machine and to write interactive notebooks

- Set it up on the remote
- Use it for the for code editing using your browser

- Connect your local installation to the remote
- Use it for the for code editing in your local PyCharm

CONCLUSIONS

CONCLUSIONS

- How to perform first steps on the server
- How to configure your environment
- How to inspect resources of the remote machine
- How to set up a sandbox
- How to write a code